

# Explainability for Deep Neural Networks

Katarína Grešová





Katarína Grešová

<https://katarinagresova.github.io>

# Outline


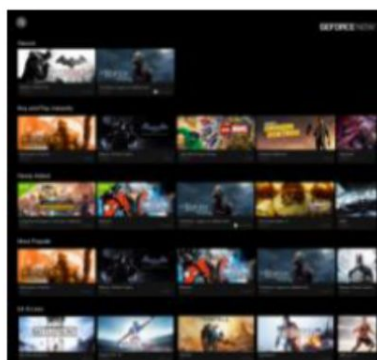

- 14:00 – 14:10 Introduction to explainability for deep neural networks
- 14:10 – 14:40 Hands on: Practical overview of explainability methods for genomic sequence data
- 14:40 – 15:30 Hands on: Practical overview of explainability methods for image data
- 15:30 – 16:00 Coffee break
- 16:00 – 16:30 Use case: miRNA target prediction
- 16:30 – 17:00 Hands on: Using DeepExperiment to interpret and visualize miRNA targeting



# Introduction to explainability for deep neural networks



# DEEP LEARNING EVERYWHERE



## INTERNET & CLOUD

Image Classification  
Speech Recognition  
Language Translation  
Language Processing  
Sentiment Analysis  
Recommendation

## MEDICINE & BIOLOGY

Cancer Cell Detection  
Diabetic Grading  
Drug Discovery

## MEDIA & ENTERTAINMENT

Video Captioning  
Video Search  
Real Time Translation

## SECURITY & DEFENSE


Face Detection  
Video Surveillance  
Satellite Imagery

## AUTONOMOUS MACHINES

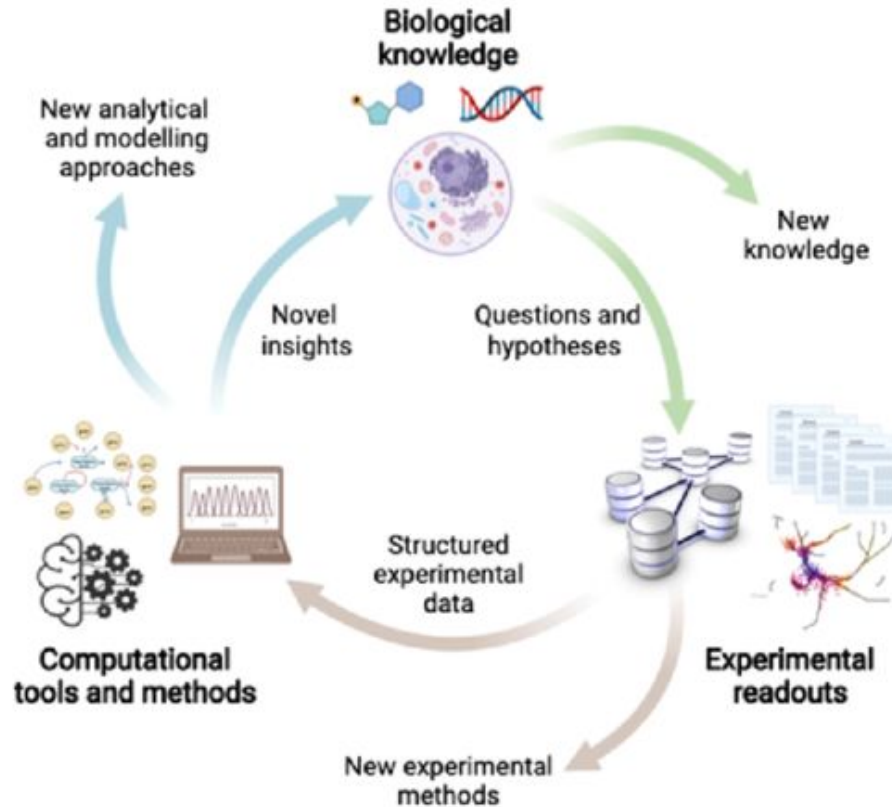
Pedestrian Detection  
Lane Tracking  
Recognize Traffic Sign

# When do we need model interpretation?


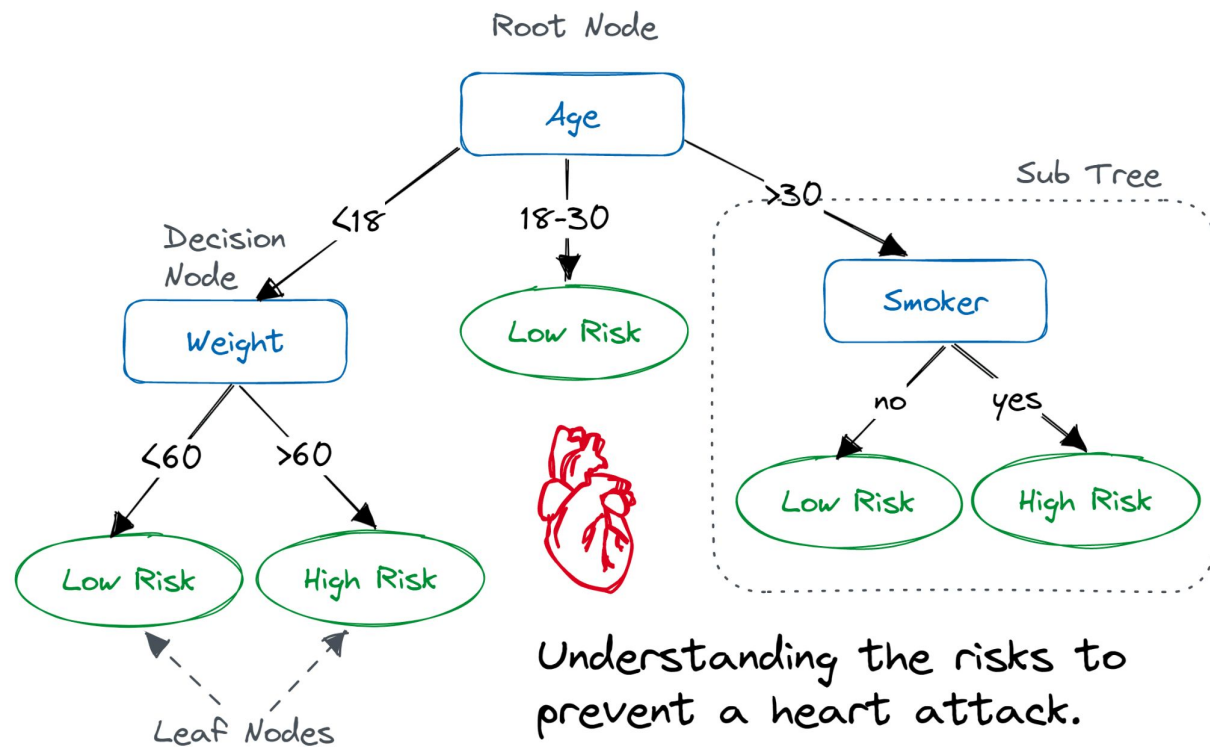
- High-stakes decision making settings
  - Impact on human lives/health/finances
  - Less studied problems, models not extensively validated
- Accuracy alone is no longer enough
  - Train/test data might not be representative of data encountered in practice
- Auxiliary criteria are also crucial
  - Nondiscrimination
  - Right to explanation



# Model interpretation as scientific method?

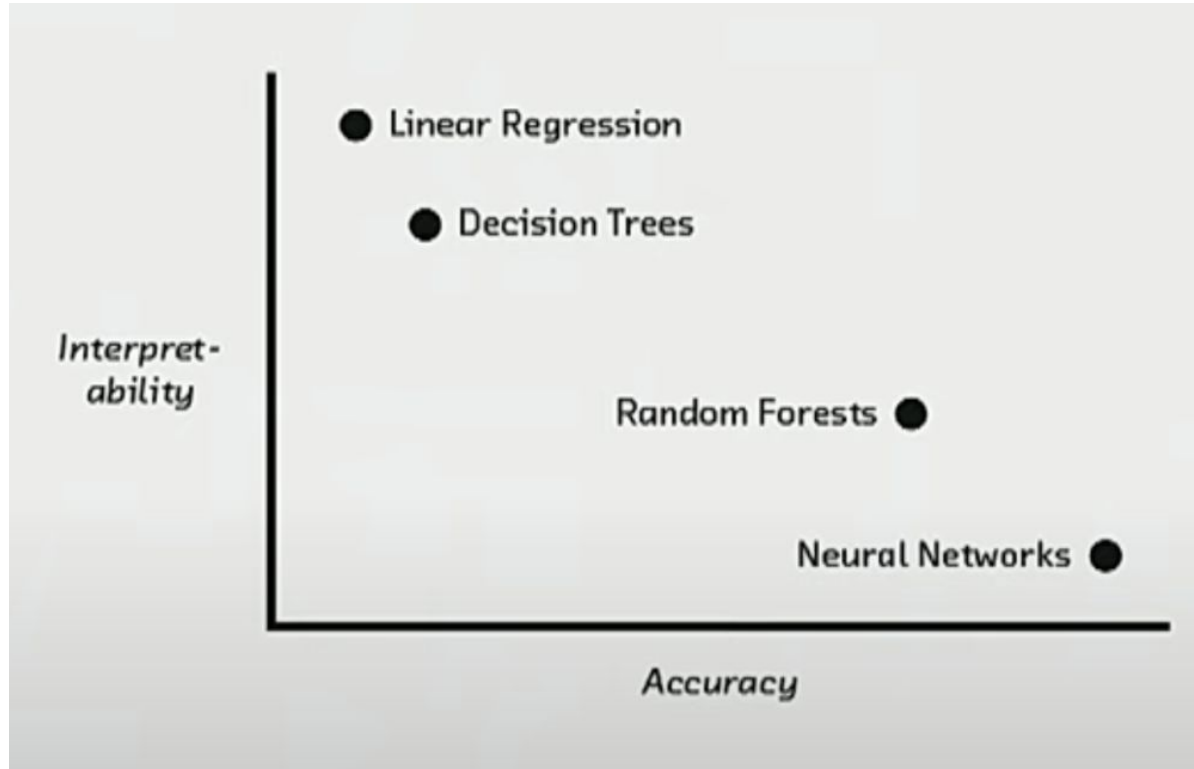


# Machine Learning vs. Deep Learning

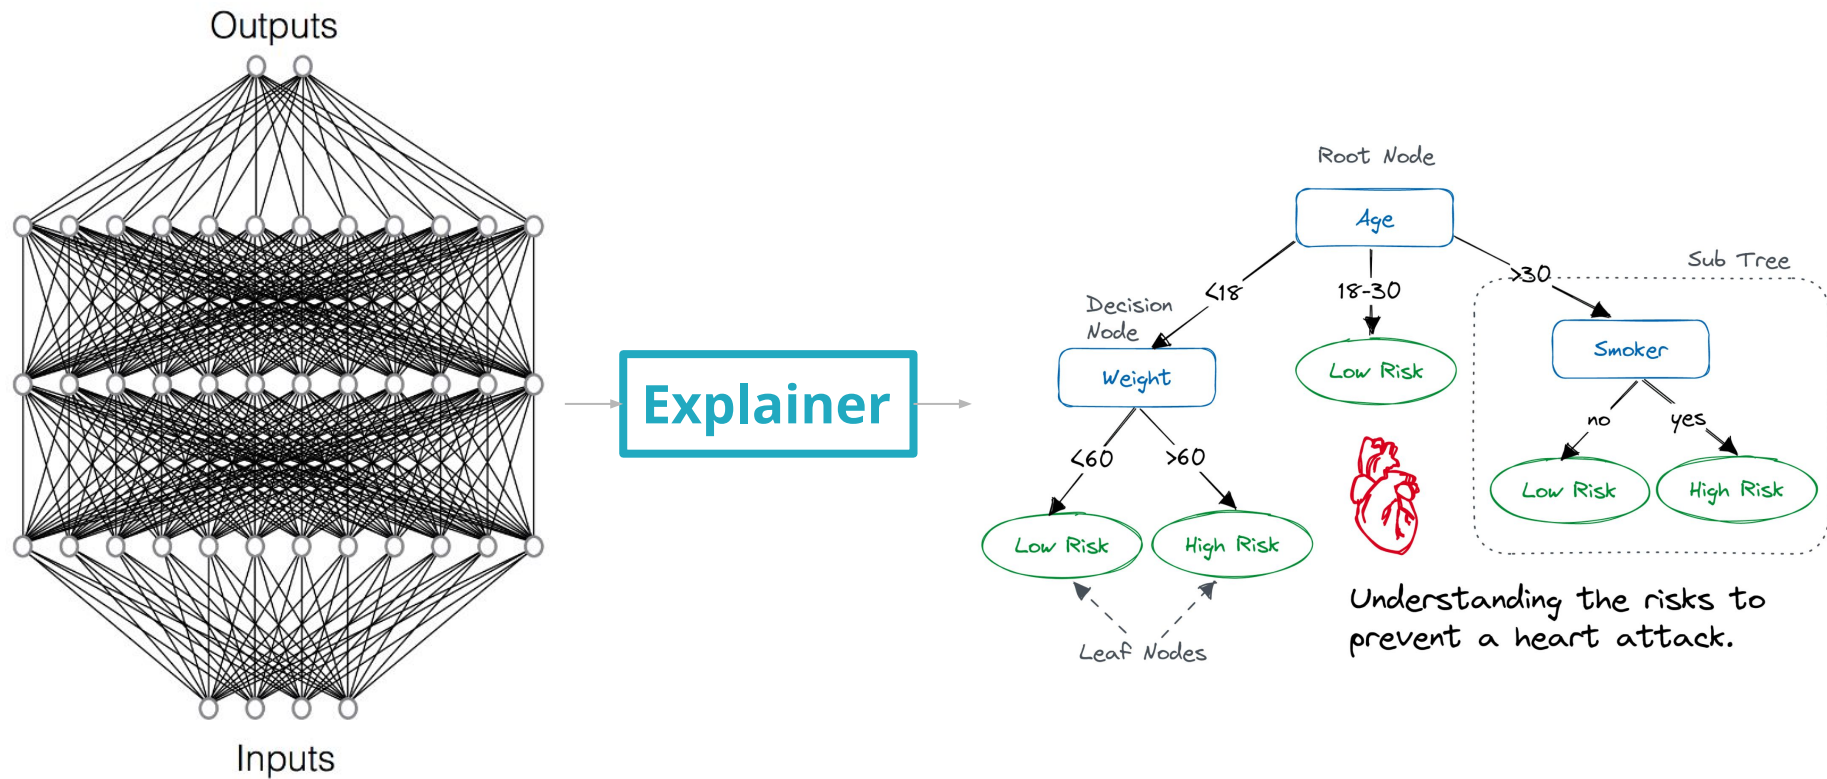




# Interpretability vs Accuracy tradeoff



# Post-hoc explainability





# Hands on: Practical overview of explainability methods for genomic sequence data



# Open the Colab notebook

Hands on: Practical overview of explainability methods for genomic sequence data (<https://colab.research.google.com/drive/1Br0f8xPIBkGFIPuXZPMtDDV8GG64wkf2?usp=sharing>)



# Hands on: Practical overview of explainability methods for image data



# Open the Colab notebook

[Hands on: Practical overview of explainability methods for image data](https://colab.research.google.com/drive/1cO54Si-hoTZkrkIRS3WYxYHVm4YQUEhB?usp=sharing)  
(<https://colab.research.google.com/drive/1cO54Si-hoTZkrkIRS3WYxYHVm4YQUEhB?usp=sharing>)



# Coffee Break!

Let's start again in 30 minutes (16:00)






Use case: miRNA target prediction







# Biological meaning




# Biological meaning




# RISC (RNA-induced silencing complex)


siRNA duplex




Ago2-RISC  
Integration



mRNA  
recognition




mRNA  
degradation




# Seed Binding + More?


siRNA duplex





Ago2-RISC  
Integration



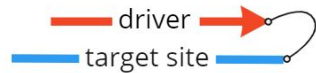


mRNA  
recognition



mRNA  
degradation

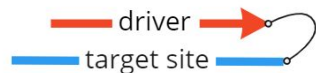




# Biological experiment - CLASH



Chimeric  
Read

# Biological experiment - CLASH

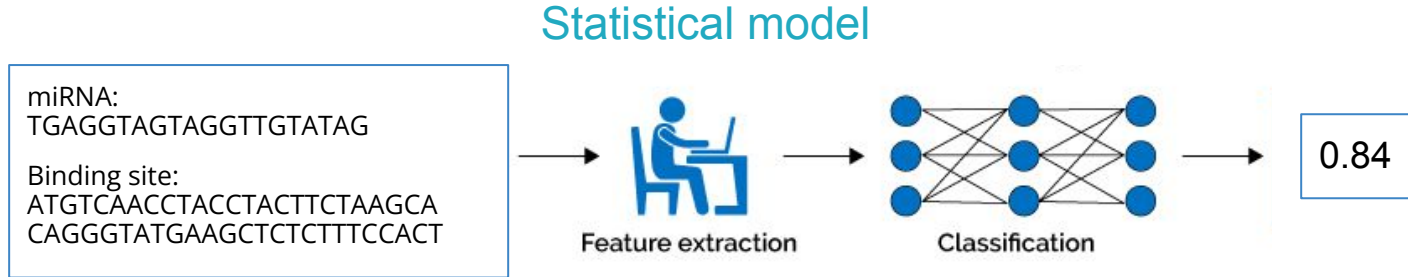


Chimeric  
Read

miRNA	gene	label
AACTGGCCCTCAAAGTCCCG	TGGAGAGCGGGCTTAAGAAGTGGCGGTTGGGCCGAGGTTCCATCGTATC	1
ATCAGGGCTTGTTGAATGGG	CTCGCTGGCGTTCTCCGGGGTGGTTGGCATTGTGTCCTGGAAGCGGCCAT	0
TGGGGAGCTGAGGCTCTGGG	CTACACCTCAGCCCGGGGCTGCACTGCCACCCTGGGCAACTTCGCCAAGG	0
GTGAGGGCATGCAGGCCTGG	GTAAGGAGCTGGAGTCGCTGGTAGAGAACGAGGGCAGTGAGGTGCTGGCG	0
ATGCACCTGGGCAAGGATTC	GCATATGGGGGCTTAAGGAATAACAGTGTGCGTGGTGGTGTGCAGGAGA	0
TGCACGGCACTGGGGACACG	TCAGGGTTTCTTGGGGGCTTATGAGTCTCACC GGTC AACC CAGGAGGCCT	0
AACTGGCCCTCAAAGTCCCG	ACCTCTTAATGGGCCAGTGAATAACACTCACTGCTGGCATTTAATGTGCA	1
TGGGTTCTTGGCATGCTGAT	CACCTGCTGCCCCCTTACCCAGCTCCACCACCTGCAGTCCCTAAAGAA	0
TCAGTGCATCACAGAACTTT	ACCCGCACAGCAAGCACCTGTACACGGCCGACATGTTACGCACGGGATC	0
CTGGCCCTCTCTGCCCTTCC	CTGATTGTGGCAGAGGGGCCACTACCCAAGGTCTAGCTAGGCCCAAGACC	1
TGAGGTAGTAGGTTGTATAG	ATGACCCAACCTACCACCCTGTTTTTACATATCCAATTCCAGTAACTCTC	1
TAAAGTGCTTATAGTGCAGG	CAAAAGCATACCTACCTTCCCCTAGAGGTCTGTAACATTGTGGCTGGGCA	1
TGAGAACTGAATTCCATGGG	CCTGGGACCCCCAGGCGTGGAGGACAGTCAAGCCGTGGAGGCCGTGGAGG	0
TGAGGTAGTAGGTTGTATAG	CCCAACCTCAACCTCAACCTCCCAGCACCACACATCATGCCAGGGGTTGG	1
CTGTACAGGCCACTGCCTTG	GAAGGTAAAGAGGGTCATTGGGGTCGAGCTATGCCCAGAGGCTGTGGAGG	0
GTCCCTCTCCAAATGTGTCT	GCTGGCCAGCGGACTTCTGGAGTTAGCCTTTGCTTTTGGAGGACTGTGTG	0
TTAGGGCCCTGGCTCCATCT	ACACAGGAAGAGGAGCCAGGCCCTTGTAACCTATGGGATTGGACAGGACTG	1
TAGGTAGTTTCATGTTGTTG	TCCGCCCTCTTTTGGCAGCCAGCCCTCCATGCACATTTGGACGCTGTC	0
TAAAGAGCCCTGTGGAGACA	TCCTGAGGCCTGGGGCACCTTTTCGTCTGATGAGCCTCTGCATGGAGAGAG	0
GTGGGTACGGCCCCAGTGGG	CATCTTGTCTCACAGCCCAGAGCATGTTCCAGATCCCAGAGTTTGAGCC	0

Helwak et al., 2013 CLASH dataset - 30 785 miRNA:target site pairs

# Computational model




# Computational model

## Statistical model




## Deep neural network







Article

## miRBind: A Deep Learning Method for miRNA Binding Classification


Eva Klimentová <sup>1,†</sup>, Václav Hejret <sup>1,2,†</sup>, Ján Krčmář <sup>3</sup>, Katarína Grešová <sup>1,2</sup> , Ilektra-Chara Giassa <sup>1,\*</sup> and Panagiotis Alexiou <sup>1</sup>



(a)




(b)



(c)


Precision-recall curves




# Interpreting Neural Networks




# Interpreting Neural Networks



GATA motif      TAL1 motif





# Interpreting Neural Networks




meerkat

mongoose



GATA motif      TAL1 motif



How to interpret  
interaction between  
sequences

TACGTCAGTTCATGAAGCT

A




(driver ~20nt)

×


AGTTCTAGTTCGTCCGTCAGTGTCAG

TTCATGAGCACCAGTCACGTTCGTCTA


(target ~50nt)



# miRBind model - interpretation




# miRBind model - interpretation




miRNA:  
TGAGGTAGTAGGTTGTATAG

Binding site:  
ATGTCAACCTACCTACTTCTAAGCA  
CAGGGTATGAAGCTCTCTTCCACT



0.84





# Visualization

miRNA: TGAGGTAGTAGGTTGTATAG

Binding site: ATGTCAACCTACCTACTTCTAAGCACAGGGTATGAAGCTCTCTTTCCACT

Predicted alignment:



# Visualization


miRNA: TGAGGTAGTAGGTTGTATAG

Binding site: ATGTCAACCTACCTACTTCTAAGCACAGGGTATGAAGCTCTCTTTCCACT

Predicted alignment:

TCACCTTTCTCTCGAAGTATGGGACACGAATCTTCATCCATCCAACGTGTA←  
-----TGAGGTA-GTAGGTTGTATAG→

miRNA position importance:



# Visualization



miRNA: TGAGGTAGTAGGTTGTATAG

Binding site: ATGTCAACCTACCTACTTCTAAGCACAGGGTATGAAGCTCTCTTTCCACT

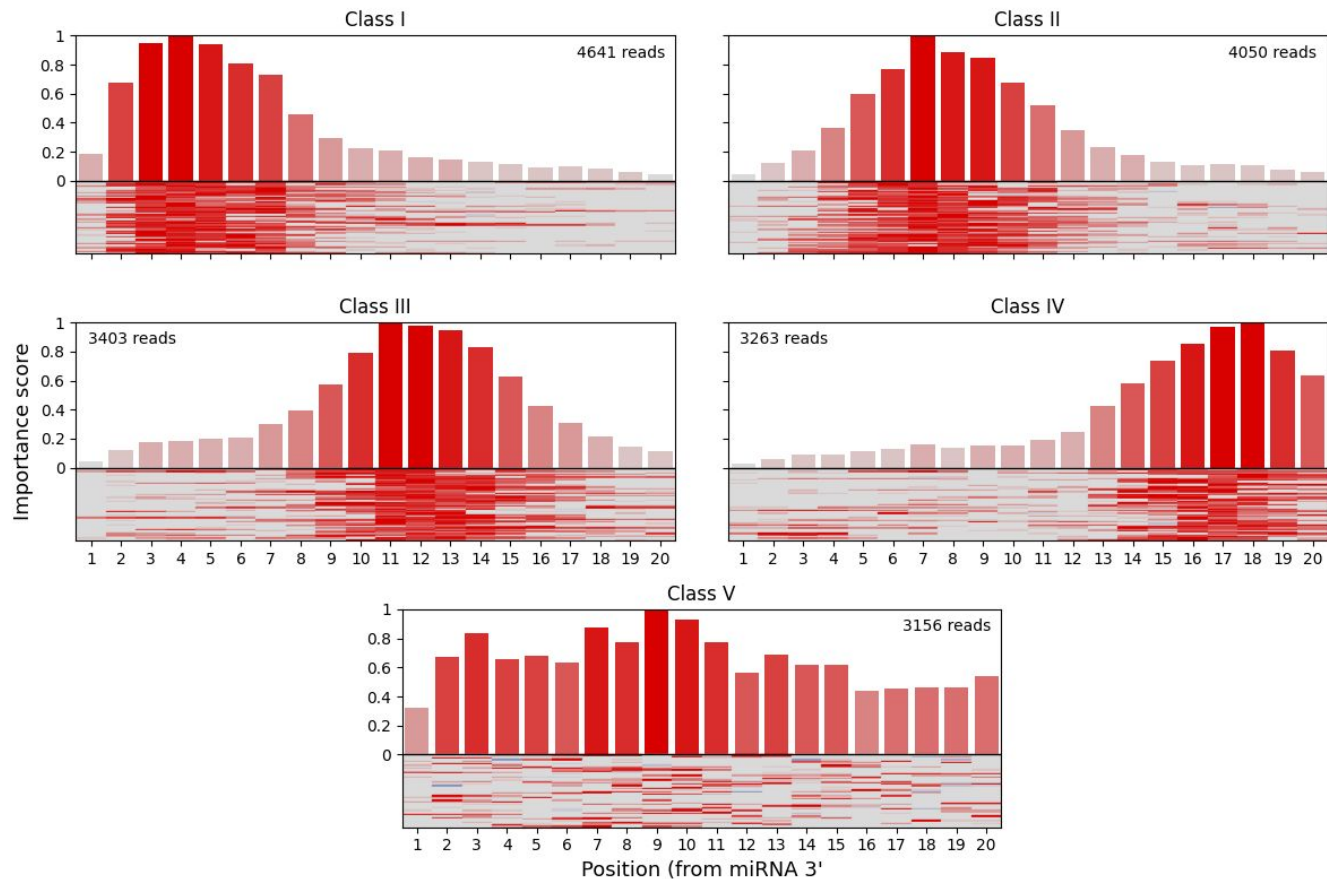
Predicted alignment:

TCACCTTTCTCTCGAAGTATGGGACACGAATCTT**CATCCATCCA**ACTGTATGTA-  
- - - - -  
- - - - -TGAGGTA-GTAGGTTGTATAG


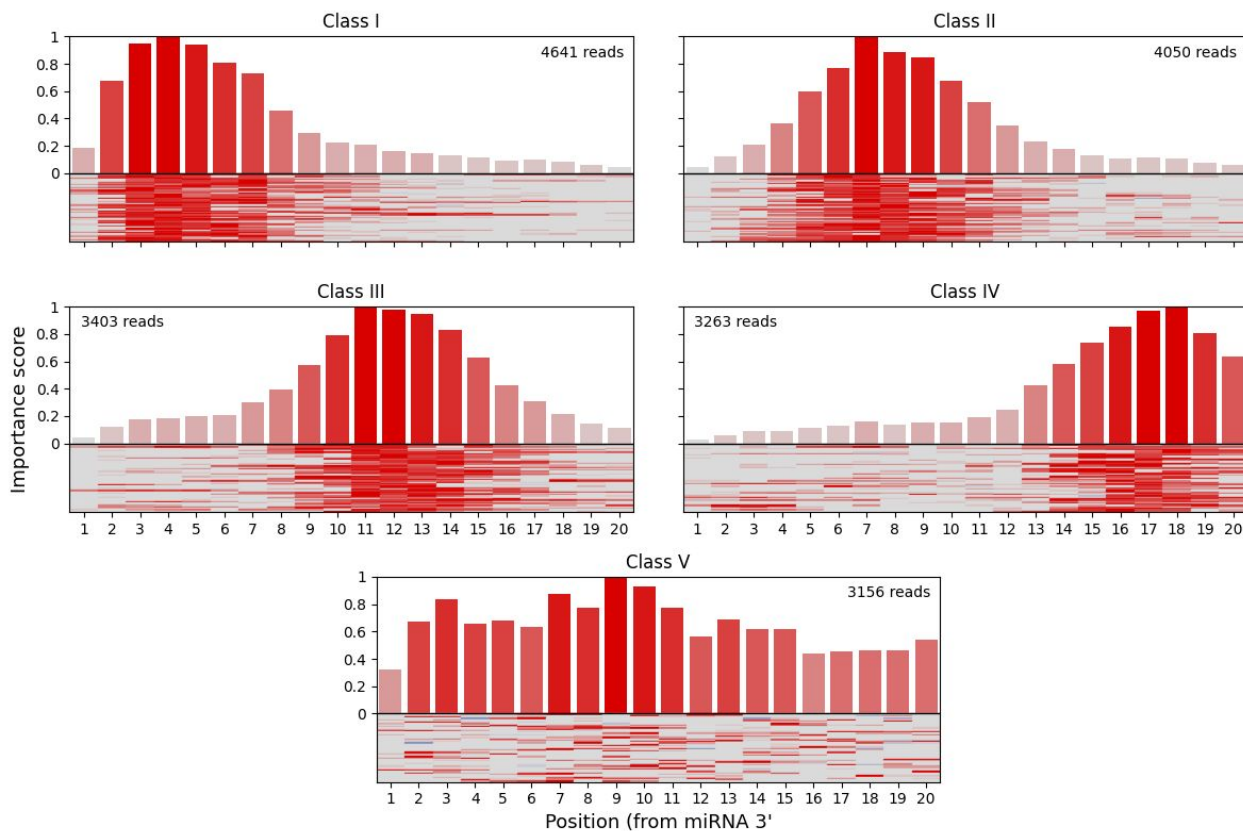
miRNA position importance:



# Classes of interaction



# Classes of interaction



Helwak et al., 2013

# Mutagenesis experiment

Open access, freely available online PLOS BIOLOGY


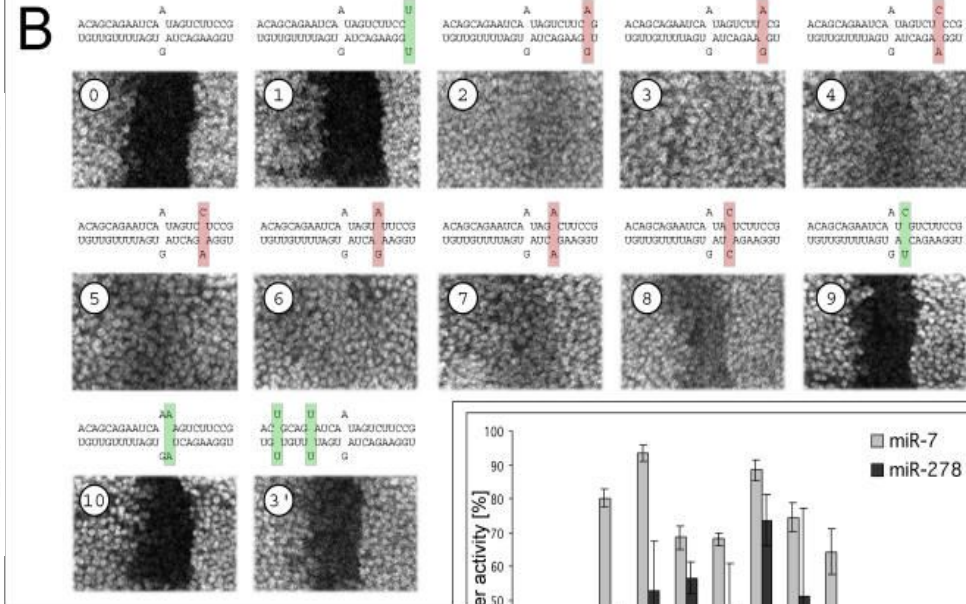

## Principles of MicroRNA–Target Recognition

Julius Brennecke<sup>1</sup>, Alexander Stark<sup>1</sup>, Robert B. Russell, Stephen M. Cohen<sup>\*</sup>

European Molecular Biology Laboratory, Heidelberg, Germany

MicroRNAs (miRNAs) are short non-coding RNAs that regulate gene expression in plants and animals. Although their biological importance has become clear, how they recognize and regulate target genes remains less well understood. Here, we systematically evaluate the minimal requirements for functional miRNA–target duplexes in vivo and distinguish classes of target sites with different functional properties. Target sites can be grouped into two broad categories. 5' dominant sites have sufficient complementarity to the miRNA 5' end to function with little or no support from pairing to the miRNA 3' end. Indeed, sites with 3' pairing below the random noise level are functional given a strong 5' end. In contrast, 3' compensatory sites have insufficient 5' pairing and require strong 3' pairing for function. We present examples and genome-wide statistical support to show that both classes of sites are used in biologically relevant genes. We provide evidence that an average miRNA has approximately 100 target sites, indicating that miRNAs regulate a large fraction of protein-coding genes and that miRNA 3' ends are key determinants of target specificity within miRNA families.

Citation: Brennecke J, Stark A, Russell RB, Cohen SM (2005) Principles of microRNA–target recognition. PLoS Biol 3(3): e85.





# Mutagenesis experiment

Open access, freely available online PLOS BIOLOGY





## Principles of MicroRNA–Target Recognition

Julius Brennecke<sup>1</sup>, Alexander Stark<sup>1</sup>, Robert B. Russell, Stephen M. Cohen<sup>\*</sup>


European Molecular Biology Laboratory, Heidelberg, Germany

MicroRNAs (miRNAs) are short non-coding RNAs that regulate gene expression in plants and animals. Although their biological importance has become clear, how they recognize and regulate target genes remains less well understood. Here, we systematically evaluate the minimal requirements for functional miRNA–target duplexes in vivo and distinguish classes of target sites with different functional properties. Target sites can be grouped into two broad categories. 5' dominant sites have sufficient complementarity to the miRNA 5' end to function with little or no support from pairing to the miRNA 3' end. Indeed, sites with strong 5' end. In contrast, 3' compensatory sites have strong 3' end. We present examples and genome-wide statistical analysis of target sites. We provide evidence that an average miRNA regulates a large fraction of protein-coding genes within miRNA families.

Citation: Brennecke J, Stark A, Russell RB, Cohen SM (2005) Principles of



# Verification



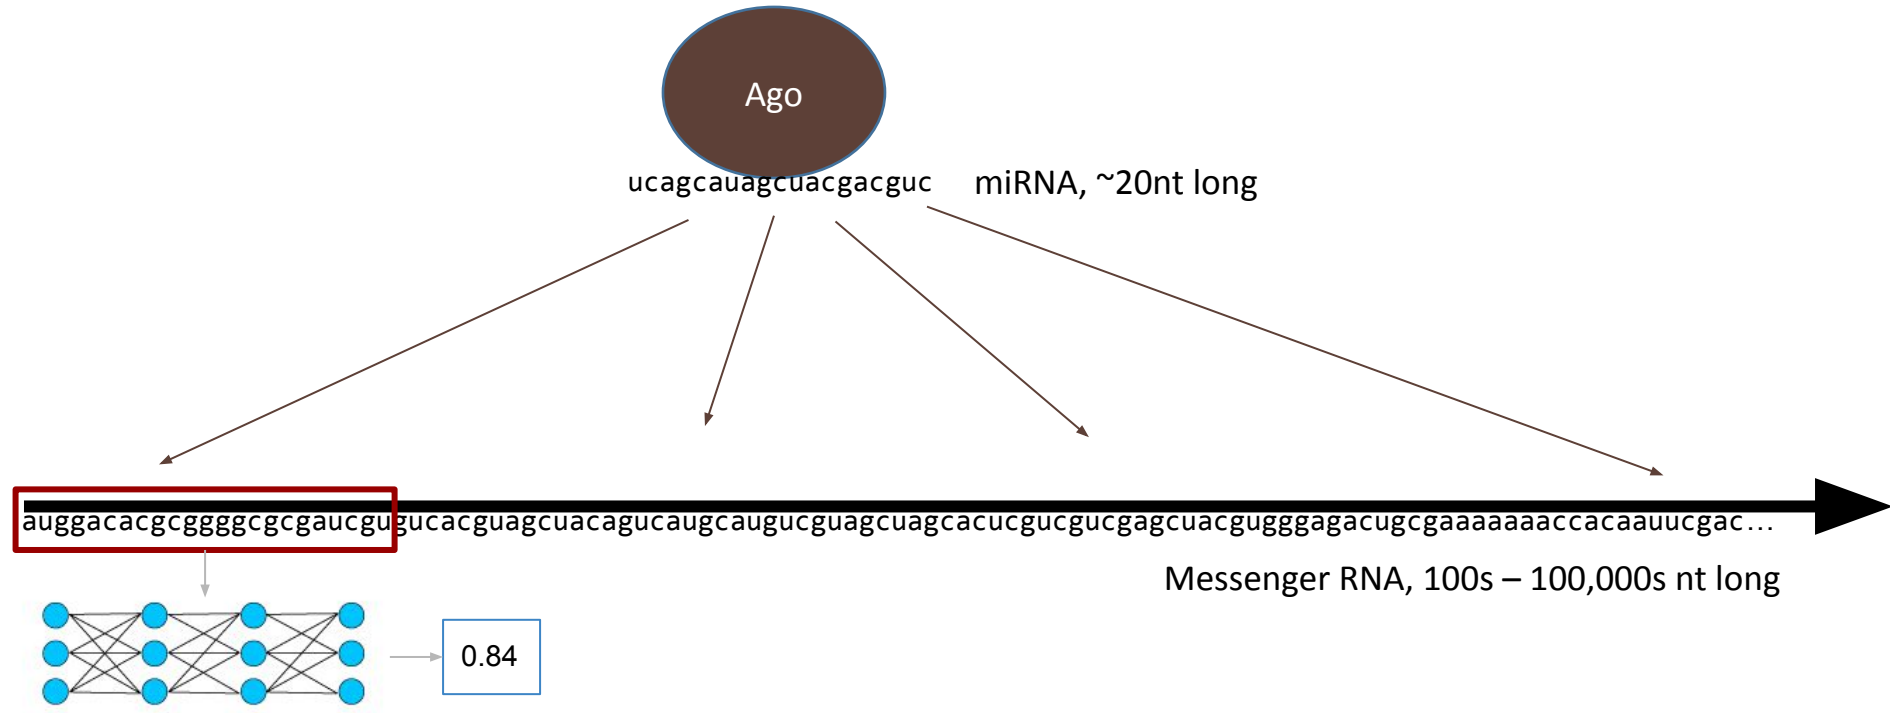
miRNA	miR-7	miR-278
correlation	0.59	0.85



# Scanning



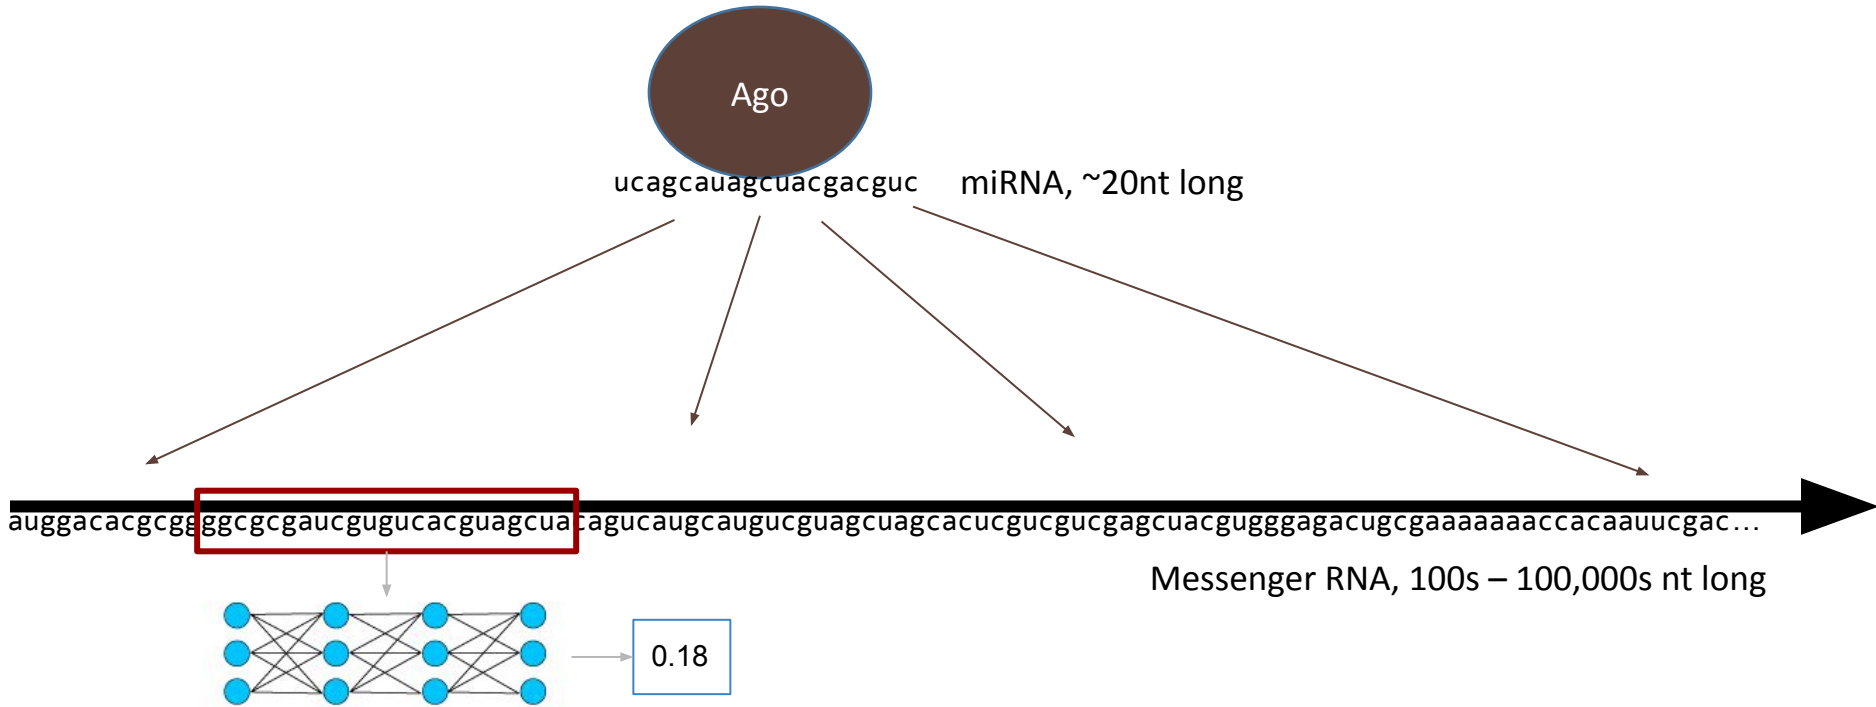
# Scanning



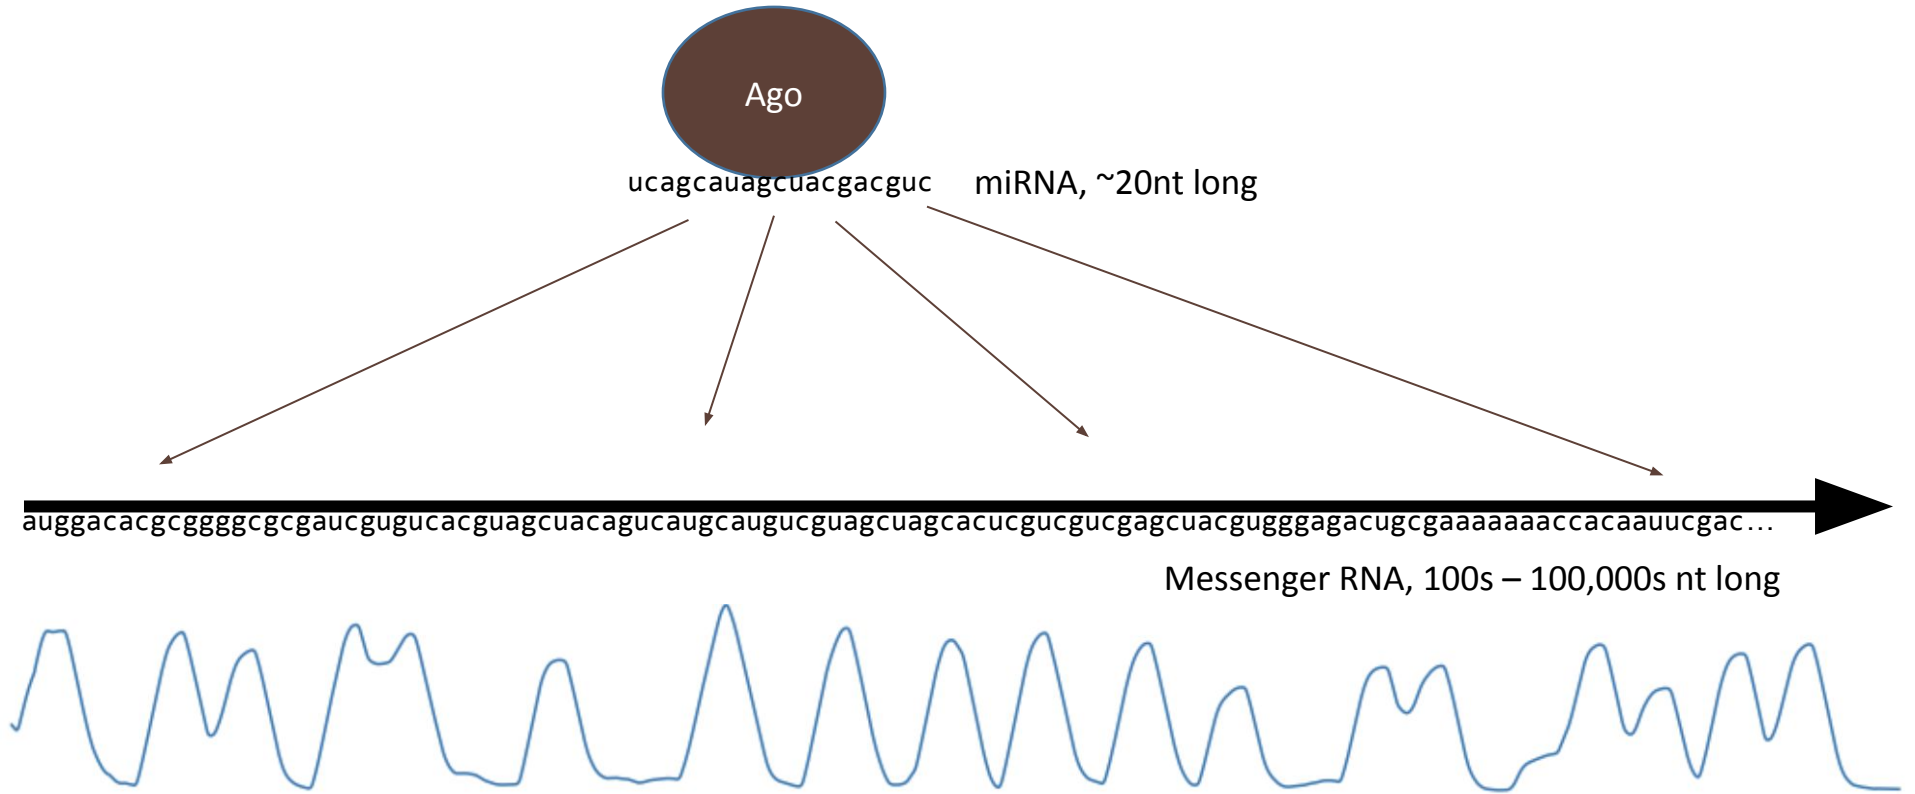
# Scanning




# Scanning




# Scanning




# Narrowing the peaks



# Narrowing the peaks



# Close by peaks







# Hands on: Using DeepExperiment to interpret and visualize miRNA targeting



# Open the Colab notebook

[Hands on: Using DeepExperiment to interpret and visualize miRNA targeting](https://colab.research.google.com/drive/1lelArVN_BJ4P9Uex3yhB8hM3MfEPGay2?usp=sharing)  
([https://colab.research.google.com/drive/1lelArVN\\_BJ4P9Uex3yhB8hM3MfEPGay2?usp=sharing](https://colab.research.google.com/drive/1lelArVN_BJ4P9Uex3yhB8hM3MfEPGay2?usp=sharing))

# Conclusions

- There are many techniques for interpreting neural networks
- They use different principles and produce different results
- Personal tip: don't use just one, try multiple interpretation techniques



# Thank you for your attention!

Deep Neural Networks are like a complex organisms and interpretation techniques help us perform experiments to better understand them.

<https://katarinagresova.github.io>